

ENGSCI 233 - Performance

Lecture 1

Colin Simpson

Semester One, 2022

Department of Engineering Science
University of Auckland
(Beamer Theme: Metropolis)

Introduction

Introduction to Performance

Performance in computational techniques relates to algorithm efficiency and how well it scales to larger problems.

We will cover several methods for identifying and improving algorithm performance:

- Measuring and comparing algorithm performance.
- Additional considerations for choosing an algorithm.
- Using a profiler to identify and fix bottlenecks.
- Implement parallelisation of code on multiple cores.

Measuring and Comparing Algorithms

Comparing Algorithms

A number of different algorithms are often available to handle a specific problem. For example, there are many different algorithms for sorting an array of numbers into ascending order.

The number of operations required to complete a problem is a useful metric for quantifying and comparing algorithm performance.

Why not just use algorithm run time to measure performance?

Algorithm Considerations

Consider the following four sorting algorithm properties:

1. Performance: number of operations required.
2. Memory: may require storing copy/copies of input data.
3. Stability (sorting): preserves original order of equal values.
4. Online: can work on input data as it is received.

Choosing a specific algorithm (e.g. for sorting) is therefore problem and hardware-dependent.


Big O Notation

Big O notation can be used to classify algorithm performance in terms of number of operations required:

- Constant e.g. $\mathcal{O}(1)$
- Logarithmic e.g. $\mathcal{O}(\log n)$
- Linear e.g. $\mathcal{O}(n)$
- Superlinear e.g. $\mathcal{O}(n \log n)$
- Polynomial e.g. $\mathcal{O}(n^2)$
- Exponential e.g. $\mathcal{O}(2^n)$
- Factorial e.g. $\mathcal{O}(n!)$

where n is a representation of the problem size (e.g. number of items in list to be sorted).

Comparison of Algorithm Performance



Best, Worst and Average Case Performance

Algorithm performance can depend on the problem itself. For example, if sorting numbers into ascending order, an algorithm may take fewer operations if the numbers start off in ascending order, and more operations if they start off in descending order.

We therefore generally consider three categories of performance:

- Best case: smallest possible number of operations.
- Worst case: largest possible number of operations.
- Average case: typical number of operations.

Sorting Algorithm Performance

Let's compare the performance of some sorting algorithms:

Algorithm	Average	Best	Worst
Heapsort	$\mathcal{O}(n \log n)$	$\mathcal{O}(n \log n)$	$\mathcal{O}(n \log n)$
Insertion sort	$\mathcal{O}(n^2)$	$\mathcal{O}(n)$	$\mathcal{O}(n^2)$
Bubble sort	$\mathcal{O}(n^2)$	$\mathcal{O}(n^2)$	$\mathcal{O}(n^2)$
Quicksort	$\mathcal{O}(n \log n)$	$\mathcal{O}(n \log n)$	$\mathcal{O}(n^2)$

Which performance category should we decide based on?

Algorithm Run Time

Big O notation gives us a sense of how algorithm performance trends as problem size grows. However, what about actual performance for a less than massive problem?


Comparing the run time, $T(n)$, of average-case heapsort and insertion sort for a less than massive problem size:

- Heapsort: $T(n) = k_1 n \log(n) + k_2$
- Insertion sort: $T(n) = k_3 n^2 + k_4$

where k_1, k_2, k_3, k_4 are problem-dependent constants. The relative values of these will determine which algorithm has better run time.

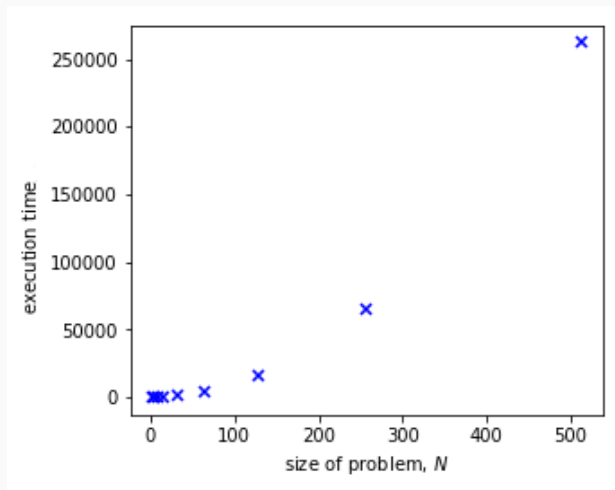
Algorithm Run Time

Comparing heapsort and insertion sort run time for varying problem size (which one is which?):




Polynomial Algorithm and Logarithmic Plots

Consider trying to determine α for a polynomial algorithm, $\mathcal{O}(n^\alpha)$.
Let's examine a plot of run time vs problem size:



Polynomial Algorithm and Logarithmic Plots

Recall the log rule $\log_b(x^y) = y \log_b(x)$. We can plot our run time and problem size on a log-log plot:



The slope of this plot therefore gives the value of α in $\mathcal{O}(n^\alpha)$.

Additional Algorithm Considerations

Other Algorithm Metrics

Number of operations is only one aspect of choosing an algorithm:

- In-Place. *Does the algorithm requires storage of many temporary values?*
- Stability. *Does the algorithm preserve order of equal numbers during sorting?*
- Online. *Can the algorithm start working on data as it is received, or does it require the full dataset?*
- Preconditions of algorithm input. *Does the algorithm only work well with certain input data?*

In-Place vs Out-Of-Place

An in-place algorithm is one that uses a constant amount of extra space, regardless of the problem size.

Insertion sort uses a single extra value to store the key, regardless of list length. The sorting process itself overwrites the existing list values. Therefore the algorithm is *in-place*.

Merge sort uses a divide and conquer approach that iteratively divides the list into halves, sorts and then merges. However, in order to do so it requires extra space that scales with list length as $\mathcal{O}(n)$. Therefore the algorithm is *out-of-place*.

Stability in Sorting

Stability here is specifically in reference to preserving the original order of equal-valued list items in a sorting algorithm.

Insertion sort will preserve order, as it iteratively inserts each item into the correct place. Therefore, the algorithm is stable.

Heapsort is not stable, as the binary heap building process removes any sense of ordering. So list items of equal value can readily swap positions in the sorted array. Therefore, the algorithm is not stable.

Online vs Offline

Some algorithms can start working with only partially complete input data.

Insertion sort takes each new list item and inserts into the currently sorted list. It can therefore continue sorting while receiving new list items, so the algorithm is *online*.

Heapsort constructs a binary heap prior to sorting. This needs to be done with the full list, so the algorithm is *offline*.

Average, Best and Worst Case

Some algorithms will work similarly regardless of the input data.

Bubble sort, which just iteratively compares and swaps adjacent elements, will take the same amount of operations regardless of the input. It has the same average, best and worst case.

Quicksort, a divide and conquer sorting algorithm, has great average and best case performance, $\mathcal{O}(n \log n)$. However, it can suffer from significantly worse worst case performance, $\mathcal{O}(n^2)$.

Learning Outcomes

Learning Outcomes

- Know what kinds of considerations go into selecting one algorithm over another.
- Interpret big O notation and its relation to performance.
- Understand the differences between average, best and worse case performance.
- Understand the differences between performance categorisation and run time.
- Calculate α in $\mathcal{O}(n^\alpha)$ from log-log plot.
- Know the meaning of in-place, stability, and online in relation to an algorithm.